



Audittens

---

# ZKsync Gateway Audit

---

December 19, 2024

# Contents

<b>1</b>	<b>About Audittens</b>	<b>4</b>
<b>2</b>	<b>About ZKsync</b>	<b>4</b>
<b>3</b>	<b>Risk classification</b>	<b>4</b>
<b>4</b>	<b>Executive summary</b>	<b>5</b>
4.1	Engagement overview . . . . .	5
4.2	Scope . . . . .	5
4.3	Summary of findings . . . . .	6
<b>5</b>	<b>Assumptions and limitations</b>	<b>7</b>
5.1	Explicit invariants of the codebase usage . . . . .	7
5.2	Limitations . . . . .	7
<b>6</b>	<b>Findings</b>	<b>8</b>
6.1	Critical severity findings . . . . .	8
6.1.1	C-01. When chain migrates to settlement layer, its priority queue there is empty but <code>priorityTree.startIndex</code> is non-zero . . . . .	8
6.2	High severity findings . . . . .	9
6.2.1	H-01. Some earliest transactions from priority tree will be executed twice . . . . .	9
6.2.2	H-02. Failed transfer recovering is possible through different asset ids . . . . .	9
6.2.3	H-03. <code>Bridgehub.setLegacyChainAddress</code> doesn't register legacy chain completely . . . . .	10
6.2.4	H-04. Failed ETH deposits with legacy <code>_data</code> are not recoverable . . . . .	11
6.2.5	H-05. <code>MailboxFacet._proveL2LeafInclusion</code> doesn't allow to prove log if the chain has exactly one batch on settlement layer . . . . .	11
6.2.6	H-06. <code>NTV.assetId</code> records can be overwritten . . . . .	12
6.2.7	H-07. Legacy withdrawals on L2 can be exploited to steal users' funds . . . . .	13
6.3	Medium severity findings . . . . .	14
6.3.1	M-01. <code>WETH</code> withdrawing and recovering is impossible . . . . .	14
6.3.2	M-02. Depositing through <code>L1ERC20Bridge</code> can spend funds twice . . . . .	14
6.3.3	M-03. <code>L2AssetRouter._ensureTokenRegisteredWithNTV</code> always returns 0 . . . . .	15
6.4	Low severity findings . . . . .	15
6.4.1	L-01. Whitelisting of settlement layers affects <code>MailboxFacet._proveL2LeafInclusion</code> of all chains . . . . .	15
6.4.2	L-02. <code>L1Nullifier._isPreSharedBridgeDepositOnEra</code> behaviour is changed from on-chain version . . . . .	15
6.4.3	L-03. Allow to call <code>AdminFacet.setPubdataPricingMode</code> even after the first batch . . . . .	16
6.4.4	L-04. In <code>RollupL1DAValidator</code> , blob can be published long before it is really used . . . . .	16
6.4.5	L-05. <code>AdminFacet.forwardedBridgeRecoverFailedTransfer</code> has redundant check . . . . .	16
6.4.6	L-06. Bridging system only partially supports L2-value . . . . .	17
6.4.7	L-07. <code>PermanentRestriction.allowL2Admin</code> is not resistant to <code>CREATE2</code> collision . . . . .	17
6.4.8	L-08. <code>L2WrappedBaseToken.initializeV2</code> fails for already deployed contract . . . . .	18
6.4.9	L-09. <code>L1AssetRouter</code> and <code>L1Nullifier</code> don't have some legacy functions for backward compatibility . . . . .	18
6.4.10	L-10. Settlement layer's admin can break system invariant . . . . .	18
6.4.11	L-11. Operator can force user to execute the same transaction twice . . . . .	19
6.5	Informational findings . . . . .	19
6.5.1	I-01. <code>ChainTypeManager.revertBatches</code> fails unless <code>ChainTypeManager</code> is manually set as validator . . . . .	19
6.5.2	I-02. <code>ChainTypeManager.registerSettlementLayer</code> always fails . . . . .	19
6.5.3	I-03. <code>MailboxFacet.proveL1ToL2TransactionStatusViaGateway</code> function is empty . . . . .	20
6.5.4	I-04. Wrong expected message length in <code>L1Nullifier._parseL2WithdrawalMessage</code> . . . . .	20
6.5.5	I-05. <code>NativeTokenVault.bridgeMint</code> emits event <code>BridgeMint</code> twice . . . . .	20



6.5.6	I-06. ERC20 getters of <code>L2BaseToken</code> always return information about Ether . . . . .	21
6.5.7	I-07. <code>IERC20.approve</code> doesn't work with certain tokens . . . . .	21
6.5.8	I-08. <code>AdminFacet.setPriorityTxMaxGasLimit</code> and <code>AdminFacet.setTokenMultiplier</code> should be <code>onlyL1</code> . . . . .	21
6.5.9	I-09. Unknown function selector is used in <code>isTransactionAllowed</code> . . . . .	21
6.5.10	I-10. Approve in <code>L1AssetRouter.depositLegacyErc20Bridge</code> is redundant . . . . .	22
6.5.11	I-11. Several functions accept <code>msg.value</code> but don't use it . . . . .	22
6.5.12	I-12. <code>L2AssetRouter.withdraw</code> and <code>L2AssetRouter.withdrawToken</code> are marked as legacy . . . . .	23
6.5.13	I-13. <code>AccessControlRestriction.validateCall</code> may miss fallback role checks for invalid selectors . . . . .	23
6.5.14	I-14. <code>L2AssetRouter.withdrawToken</code> supports both native tokens and tokens that came from another L2 . . . . .	23
6.5.15	I-15. <code>PermanentRestriction.validateRemoveRestriction</code> blocks self-calls with data shorter than 4 bytes . . . . .	24
6.5.16	I-16. <code>BridgeHelper.getERC20Getters</code> doesn't check whether calls are successful . . . . .	24
6.5.17	I-17. <code>GettersFacet.isFunctionFreezable</code> and <code>GettersFacet.isFacetFreezable</code> are inconsistent . . . . .	24
6.5.18	I-18. New L3 → L1 message proof is not completely backward-compatible . . . . .	24
6.5.19	I-19. Priority operation expiration timestamp is not stored in priority tree . . . . .	25
6.5.20	I-20. Priority operation expiration should be increased for chains on settlement layer . . . . .	25
6.5.21	I-21. In <code>ChainTypeManager.deployNewChain</code> , condition " <code>getZKChain(_chainId) != address(0)</code> " is never satisfied . . . . .	25
6.5.22	I-22. System design doesn't allow to register custom asset handlers for L2-native assets . . . . .	26
6.5.23	I-23. Admin of a legacy chain can make <code>baseTokenAssetId</code> registered in <code>Bridgehub</code> . . . . .	26
6.5.24	I-24. Operator can force <code>Executor</code> to store <code>systemLogs</code> inconsistent with their hash . . . . .	27
6.5.25	I-25. Factory deps are not checked during the creation of new chain . . . . .	27
6.5.26	I-26. <code>AssetRouterBase</code> uses an incorrect storage gap size . . . . .	27
6.5.27	I-27. Misleading comment in <code>ChainTypeManager.registerSettlementLayer</code> regarding settlement layer deployment . . . . .	28
6.5.28	I-28. Misleading use of <code>IBridgehub</code> in <code>abi.encodeCall</code> . . . . .	28
6.5.29	I-29. Misleading variable names corresponding to <code>bridgehubData</code> in <code>Bridgehub</code> contract . . . . .	28
6.5.30	I-30. Redundant check for <code>ASSET_ROUTER</code> in <code>L1NativeTokenVault.receive</code> . . . . .	28
6.5.31	I-31. Usage of raw <code>keccak256</code> computation instead of <code>DataEncoding.encodeAssetId</code> . . . . .	29
6.5.32	I-32. Incorrect comment in <code>NativeTokenVault.bridgeBurnNativeToken</code> function regarding <code>msg.value</code> check . . . . .	29
6.5.33	I-33. Inconsistency between <code>MessageRoot</code> implementation and documentation . . . . .	29
6.5.34	I-34. Incorrect implementation in the <code>L2AssetRouter's onlyAssetRouterCounterpartOrSelf</code> modifier . . . . .	29
6.5.35	I-35. Inconsistent parameter naming of <code>_l1Asset</code> and <code>_l1Token</code> . . . . .	30
6.5.36	I-36. Inconsistent use of <code>SUPPORTED_ENCODING_VERSION</code> constant in <code>BatchDecoder</code> . . . . .	30
6.5.37	I-37. Inconsistent variable naming for chain identifiers . . . . .	30
6.5.38	I-38. Inconsistent naming of <code>assetAddress</code> and <code>assetHandlerAddress</code> in functions and events . . . . .	30
6.5.39	I-39. Inconsistent behaviour of the legacy <code>L1AssetRouter.finalizeWithdrawal</code> function between legacy and new chains . . . . .	31
6.5.40	I-40. Lack of reentrancy control in <code>L2AssetRouter</code> and <code>L2SharedBridgeLegacy</code> . . . . .	31



# 1 About Audittens

Audittens is an audit company with expertise across various sectors of Web3 security.

What truly sets us apart is that our team consists entirely of top participants in mathematics and competitive programming competitions, uniquely equipping us to rapidly analyze codebases, uncover explicit and implicit invariants, and identify unique attack vectors.

You can subscribe and learn more about us at <https://x.com/Audittens>.

# 2 About ZKsync

ZKsync is a Layer-2 protocol that scales Ethereum with cutting-edge ZK tech. Their mission is not only to merely increase Ethereum's throughput, but to fully preserve its foundational values — freedom, self-sovereignty, decentralization — at scale.

ZK-rollups like Era are the only scaling solution that can inherit 100% of Ethereum's security. But theory is not enough. ZKsync is committed to go above and beyond to make Era by far the most secure L2, in practice.

You can subscribe and learn more about the project at <https://zksync.io>.

# 3 Risk classification

The severities of the issues are determined based on the following properties:

- Critical — results in a significant loss of assets within the protocol, a major deviation from the expected behavior of protocol components, and/or a violation of key security invariants.
- High — causes loss of assets within the protocol and/or general violations of protocol security invariants, with limitations on the variety of potential attacks.
- Medium — enables barely profitable attacks on the protocol, violations of security invariants that pose minimal risk to protocol users, and/or functionality limitations affecting only a relatively small subset of users.
- Low — enables griefing attacks on the protocol, unexpected changes in the protocol's behavior that are imperceptible, and/or functionality limitations with negligible impact on the security of the protocol.
- Informational — issues that have no practical impact on the security of the current version of the protocol but could potentially lead to more serious consequences in future updates, create the possibility for incorrect use of the protocol by users, and/or result in a significant decline in code quality, making maintenance more challenging.

While determining the severity of issues, the constraints required for successful attacks and the potential actions to address their consequences are taken into account in the decision-making process.



## 4 Executive summary

### 4.1 Engagement overview

ZKsync Security Council Foundation engaged Audittens to review the security of ZKsync's bridging and chain migration protocol. From September 23, 2024, to November 29, 2024, a team of five security researchers audited the provided source code. After the completion of the fix period by ZKsync's team, the full audit was finalized by reviewing the corresponding commits.

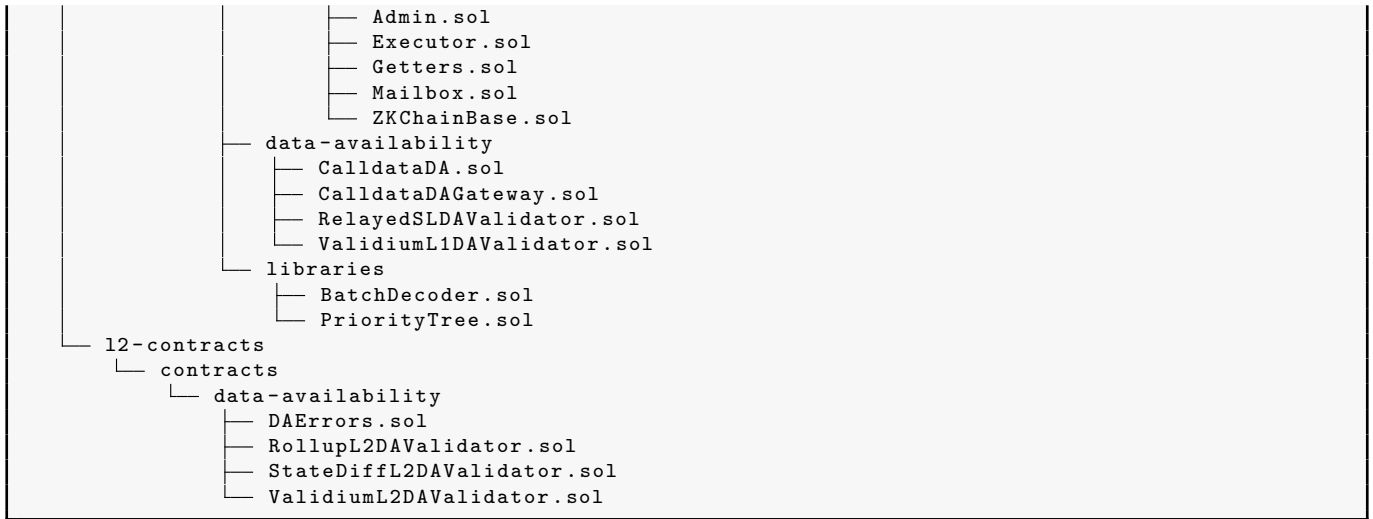
### 4.2 Scope

The `matter-labs/era-contracts` repository at the [7198b54fbcba37aa7a1dd75fc3067391af33e03e](#) commit was audited.

The following contracts were within the scope of the audit:

```
matter-labs
├── era-contracts
│   ├── da-contracts
│   │   └── contracts
│   │       ├── CalldataDA.sol
│   │       ├── DAContractsErrors.sol
│   │       └── RollupL1DAValidator.sol
│   └── l1-contracts
│       ├── contracts
│       │   ├── bridge
│       │   │   ├── BridgeHelper.sol
│       │   │   ├── BridgedStandardERC20.sol
│       │   │   ├── L1ERC20Bridge.sol
│       │   │   ├── L1Nullifier.sol
│       │   │   ├── L2SharedBridgeLegacy.sol
│       │   │   ├── L2WrappedBaseToken.sol
│       │   │   └── asset-router
│       │   │       ├── AssetRouterBase.sol
│       │   │       ├── L1AssetRouter.sol
│       │   │       └── L2AssetRouter.sol
│       │   └── ntv
│       │       ├── L1NativeTokenVault.sol
│       │       ├── L2NativeTokenVault.sol
│       │       └── NativeTokenVault.sol
│       ├── bridgehub
│       │   ├── Bridgehub.sol
│       │   ├── CTMDeploymentTracker.sol
│       │   └── MessageRoot.sol
│       ├── common
│       │   └── libraries
│       │       └── DataEncoding.sol
│       ├── governance
│       │   ├── AccessControlRestriction.sol
│       │   ├── ChainAdmin.sol
│       │   ├── Common.sol
│       │   ├── L2AdminFactory.sol
│       │   ├── L2ProxyAdminDeployer.sol
│       │   ├── PermanentRestriction.sol
│       │   ├── TransitionalOwner.sol
│       │   └── restriction
│       │       ├── Restriction.sol
│       │       └── RestrictionValidator.sol
│       ├── state-transition
│       │   ├── ChainTypeManager.sol
│       │   ├── ValidatorTimelock.sol
│       │   └── chain-deps
│       │       ├── DiamondInit.sol
│       │       ├── ZKChainStorage.sol
│       │       └── facets
```





Additionally, the differences of the files in the scope of the audit between the [7198b54f](#) and [a5754174](#) commits were separately briefly reviewed.

### 4.3 Summary of findings

Status Severity	Acknowledged	Partially fixed	Fixed	Total
Critical	0	0	1	1
High	0	0	7	7
Medium	0	0	3	3
Low	2	1	8	11
Informational	8	0	32	40
Total	10	1	51	62

Table 1. Distribution of found issues.



## 5 Assumptions and limitations

### 5.1 Explicit invariants of the codebase usage

The audit was conducted assuming that:

1. All new contracts on L1 will be deployed and configured before the upgrade of the existing contracts. In particular, `L1NativeTokenVault.registerEthToken`, `L1AssetRouter.setNativeTokenVault` and `MessageRoot.initialize` will be called before the current `Bridgehub` will be upgraded.
2. All existing contracts on L1 will be upgraded atomically in one multicall. In the same multicall all necessary initialization functions will be called (in particular, `Bridgehub.initializeV2` and `Bridgehub.setAddresses`).

### 5.2 Limitations

The codebase was audited as is. Any subsequent changes may introduce new vulnerabilities and require a separate audit. Fixes for all findings have been reviewed within the limited scope (in the context of only relevant protocol components).



## 6 Findings

### 6.1 Critical severity findings

#### 6.1.1 C-01. When chain migrates to settlement layer, its priority queue there is empty but `priorityTree.startIndex` is non-zero

**Context:**

[l1-contracts/contracts/state-transition/chain-deps/facets/Admin.sol:329](#)  
[l1-contracts/contracts/state-transition/libraries/PriorityTree.sol:87](#)  
[l1-contracts/contracts/state-transition/chain-deps/facets/Mailbox.sol:573](#)  
[l1-contracts/contracts/state-transition/chain-deps/facets/Executor.sol:444](#)  
[l1-contracts/contracts/state-transition/chain-deps/facets/Admin.sol:369](#)

**Description:** Suppose there is a chain which used priority queue, and then switched to priority tree. Its priority queue contains transactions with indices from 0 to `priorityTree.startIndex - 1` (or more), and priority tree contains transactions with indices from `priorityTree.startIndex` to `priorityTree.startIndex + priorityTree.tree._nextLeafIndex - 1`.

When this chain migrates to settlement layer, its state is copied. The priority tree is copied too, including `priorityTree.startIndex`, but priority queue on settlement layer is left empty.

Many checks in the codebase depend on `priorityQueue.getFirstUnprocessedPriorityTx` which will be equal to 0 on settlement layer. Until priority queue length reaches `priorityTree.startIndex`, the following will be happening on settlement layer:

- [Mailbox.sol:573](#) — mailbox will store transactions in both priority queue and priority tree (but these transactions have different indices);
- [Executor.sol:444](#) — executor will execute transactions from priority queue;
- [Admin.sol:369](#) — migration back will not be possible.

Therefore, transaction which is requested once on L1 will be executed twice on L2 (the first time when transaction is requested and the second time when priority queue length reaches `priorityTree.startIndex`). The attacker can abuse this by depositing assets: on L1, token will be transferred from the attacker to bridge only once but it will be minted on L2 twice (once immediately when the transaction from `priorityQueue` is processed, and once later when the same transaction from `priorityTree` is processed). The attacker can immediately withdraw asset (not waiting for the second mint) and repeat this process many times to steal all chain's balance.

Another issue is that transactions which were already present in priority tree will have to wait for a long time until priority queue length reaches `priorityTree.startIndex`. Also, when priority queue length reaches `priorityTree.startIndex`, priority tree will contain many transactions, so future transactions will have to wait too.

**Recommendation:** After line [Admin.sol:329](#), add the following statements:

```
s.priorityQueue.head = s.priorityTree.startIndex;  
s.priorityQueue.tail = s.priorityTree.startIndex;
```

**ZKsync:** *Fixed* in the [4039eab5](#) commit.

**Audittens:** *Approved*.





## 6.2 High severity findings

### 6.2.1 H-01. Some earliest transactions from priority tree will be executed twice

#### Context:

[l1-contracts/contracts/state-transition/chain-deps/facets/Mailbox.sol:572-582](#)

[l1-contracts/contracts/state-transition/chain-deps/facets/Executor.sol:443-453,387,349,419](#)

**Description:** If chain used priority queue, and then introduces priority tree, it goes into an intermediate state until all transactions requested before priority tree introducing are processed.

When user requests new priority transaction in this state, this transaction is pushed to both queue and tree.

`ExecutorFacet.executeBatchesSharedBridge` executes each batch one by one. For each batch: while chain is in intermediate state, it calls one overload of `_executeOneBatch` (which works only with `priorityQueue`); as soon as chain comes out of the intermediate state, it calls another overload of `_executeOneBatch` (which works only with `priorityTree`).

Likely, there will be also “mixed” batch which contains both some latest transactions from priority queue and some earliest transactions from priority tree. When `ExecutorFacet.executeBatchesSharedBridge` processes this batch, it will call the former overload of `_executeOneBatch` (which works only with `priorityQueue`), so all these transactions will be processed from the queue’s point of view, but none of them will be processed from tree’s point of view. When the function will process the next batch, it will call the latter overload of `_executeOneBatch` (which works only with `priorityTree`) but since some earliest transactions from priority tree are not yet processed (from the priority tree’s point of view), they’ll have to be processed again.

**Recommendation:** After line `_executeOneBatch(batchesData[i], i);`, add the statement:

```
if(s.priorityTree.startIndex < s.priorityQueue.getFirstUnprocessedPriorityTx())
    s.priorityTree.unprocessedIndex =
        s.priorityQueue.getFirstUnprocessedPriorityTx() - s.priorityTree.startIndex;
```

**ZKsync:** *Fixed* in the [519729e6](#) commit.

**Audittens:** *Approved.*

### 6.2.2 H-02. Failed transfer recovering is possible through different asset ids

#### Context:

[l1-contracts/contracts/bridge/L1Nullifier.sol:385,481](#)

[l1-contracts/contracts/common/libraries/DataEncoding.sol:108-111,134](#)

**Description:** Function `_isLegacyTxDataHash` does not check that the provided `_assetId` was used for the deposit. Instead, it checks that the corresponding `tokenAddress` from `L1NativeTokenVault` was used for the deposit. It means, that different `_assetId` that corresponds to the same `tokenAddress` can be used for the successful verification of failed transfers. It can be exploited by a malicious attacker to manipulate chain balances and permanently freeze other users’ funds.

**Attack scenario:** Imagine there is a native token on L2A (let `ID1` be the corresponding asset id for it) held by two users: first user holds 100 ETH of `ID1` and the second — 10 ETH of `ID1`. The first (regular) user does the following:

1. Withdraws 100 ETH of `ID1` to L1. This creates a bridged token on L1 (let `T1` be the address at which it’s deployed). After that, `L1NTV.originChainId[ID1] = L2A`, `L1NTV.tokenAddress[ID1] = T1`, `L1NTV.assetId[T1] = ID1`.
2. Deposits 100 ETH of `ID1` to some other chain L2B. Since the token is deposited to non-origin chain, `L1NTV.chainBalance[L2B][ID1]` is increased to 100 ETH by [NativeTokenVault.sol:209](#).

Now the second user (the attacker) does the following to permanently freeze the first user’s funds:

1. Withdraws 10 ETH of `ID1` to L1.
2. Registers `T1` by [NativeTokenVault.sol:86](#). It registers a new asset id `ID2`, assigns `L1NTV.tokenAddress[ID2] = T1` and reassigns `L1NTV.assetId[T1] = ID2`.



3. Initiates deposit of 10 ETH of T1 to L2B with provided zero gas limit at [Bridgehub.sol:557](#) and legacy-encoded data. Since `LEGACY_ENCODING_VERSION` is being used, `assetId = L1NTV.assetId[T1] = ID2` will be deposited ([L1AssetRouter.sol:253,376,387](#)). Since the token is deposited to non-origin chain, `L1NTV.chainBalance[L2B][ID2]` is increased to 10 ETH by [NativeTokenVault.sol:274](#). Legacy `txDataHash` is calculated at [L1AssetRouter.sol:271](#) and later stored in [L1Nullifier.sol:280](#).
4. Calls `bridgeRecoverFailedTransfer` with wrong `_assetId = ID1` ([L1Nullifier.sol:310](#)) after deposit fails on L2B. Since `L1NTV.tokenAddress[ID1] = L1NTV.tokenAddress[ID2]`, `_isLegacyTxDataHash` check at [L1Nullifier.sol:385](#) returns true, and L1NTV is forced to mint 10 ETH ([NativeTokenVault.sol:132](#)) of T1 to the attacker. Besides minting T1, L1NTV also decreases `L1NTV.chainBalance[L2B][ID1]` by 10 ETH.
5. Repeats steps 3-4 nine more times.

As a result of the operations above, the attacker manipulated chain balances almost at no cost (only paying for gas), making `L1NTV.chainBalance[L2B][ID1] = 0` and `L1NTV.chainBalance[L2B][ID2] = 100 ETH`, while in fact the only deposited asset to L2B is ID1. Since ID1 is a non-native token on L2B, the first regular user can try to withdraw it only with new message format ([L2AssetRouter.sol:174](#)) which includes the correct `_assetId = ID1` into withdrawal message. Therefore when finalizing withdrawal on L1, the first user can use only `assetId = ID1` as well ([L1Nullifier.sol:620](#)), but since `L1NTV.chainBalance[L2B][ID1] = 0`, he will not be able to finalize his withdrawal, resulting in permanent loss of funds (both on L2B and on L1).

Nobody, including the attacker, will be able to manipulate chain balances back in a similar way (by depositing ID1 and recovering failed transfer of ID2). It's because at step 3, `LEGACY_ENCODING_VERSION` has to be used, which allows to deposit only `assetId = L1NTV.assetId[T1] = ID2`.

**Recommendation:** It's possible to fix the issue in two different ways depending on the project's needs:

1. If the functionality of depositing non-native assets from L1 to L2 through `LEGACY_ENCODING_VERSION` should exist, it's enough to forbid native registration of the bridged tokens ([NativeTokenVault.sol:340](#)) by adding the require statement `require(assetId[_nativeToken] == 0);`. However, such an approach limits the functionality of using bridged tokens as separate native assets, and what is more important creates an implicit invariant for the `_isLegacyTxDataHash` function that one `tokenAddress` can always correspond to at most one asset id. Such an invariant should always be taken into account during future upgrades.
2. Otherwise, such deposits can be forbidden by adding the following line after [DataEncoding.sol:109](#): `require(_assetId == encodeNTVAssetId(block.chainid, tokenAddress));`. Such fix will automatically explicitly guarantee that only one exact `_assetId` corresponding to the L1-native token can successfully pass verification at `_isLegacyTxDataHash` function ([L1Nullifier.sol:481](#)).

**ZKsync:** *Fixed* in the [7673edf9](#) commit.

**Audittens:** *Approved*.

### 6.2.3 H-03. Bridgehub.setLegacyChainAddress doesn't register legacy chain completely

**Context:**

11-contracts/contracts/bridgehub/Bridgehub.sol:[229-255,384-391](#)

**Description:** To create new chain, `Bridgehub.createNewChain` does the following steps:

- set `chainTypeManager[_chainId]`,
- set `baseTokenAssetId[_chainId]`,
- set `settlementLayer[_chainId]`,
- call `CTM.createNewChain(...)`,
- set `zkChainMap[_chainId]`,
- call `messageRoot.addNewChain(_chainId)`.

If the chain is already registered before the currently reviewed version of the code is introduced, its `chainTypeManager[_chainId]` will already be set, and `CTM.createNewChain(...)` will already have been



called. To set `baseTokenAssetId[_chainId]`, anyone can call `Bridgehub.setLegacyBaseTokenAssetId`. To set `zkChainMap[_chainId]`, anyone can call `Bridgehub.setLegacyChainAddress`. However, there is no way to set `settlementLayer[_chainId]` and call `messageRoot.addNewChain(_chainId)`. Without these steps, chain cannot function properly:

- without `settlementLayer[_chainId]`, it cannot be migrated ([Bridgehub.sol:700](#));
- without `messageRoot.addNewChain(_chainId)`, new batches cannot be executed ([Executor.sol:397](#)).

**Recommendation:** Instead of separate functions `Bridgehub.setLegacyBaseTokenAssetId` and `Bridgehub.setLegacyChainAddress`, it would be easier to create a single function `Bridgehub.registerLegacyChain` which does all steps together:

- set `baseTokenAssetId[_chainId]`,
- set `settlementLayer[_chainId]`,
- set `zkChainMap[_chainId]`,
- call `messageRoot.addNewChain(_chainId)`.

**ZKsync:** *Fixed* in the [a11df6a1](#) commit.

**Audittens:** *Approved*.

#### 6.2.4 H-04. Failed ETH deposits with legacy `_data` are not recoverable

**Context:**

[l1-contracts/contracts/bridge/asset-router/L1AssetRouter.sol:276](#)

[l1-contracts/contracts/bridge/ntv/NativeTokenVault.sol:259-263](#)

**Description:** Currently deployed `L1SharedBridge` requires a user to provide `_depositAmount == 0` to deposit `msg.value` into the non-ETH based chain. For backward compatibility, zero `_depositAmount` is also accepted by the new [NativeTokenVault.sol:262](#) and is substituted with the `msg.value`. However, when calculating `txDataHash`, `_depositAmount` is not being substituted with the correct value by [L1AssetRouter.sol:271-277](#).

Therefore, in case such ETH deposit to the non-ETH-based chain fails on L2, it will be impossible to use the recovery mechanism to restore the funds on L1.

**Recommendation:** The easiest option would be to remove the reassignment of `_depositAmount` in the [NativeTokenVault.sol:259-263](#). However, if backward compatibility has to be preserved, extra logic should be added to the [L1AssetRouter.sol](#) while calculating the `txDataHash`. For this `transferData` has to be additionally parsed and replaced with the correct data for the special case of depositing ETH through the `NativeTokenVault`.

**ZKsync:** *Fixed* in the [4f1dde90](#) commit.

**Audittens:** *Approved*.

#### 6.2.5 H-05. MailboxFacet.`_proveL2LeafInclusion` doesn't allow to prove log if the chain has exactly one batch on settlement layer

**Context:**

[l1-contracts/contracts/state-transition/chain-deps/facets/Mailbox.sol:226-228,247](#)

[l1-contracts/contracts/common/libraries/Merkle.sol:52,121-123](#)

**Description:** In `MessageRoot`, each chain is represented as a root of a Merkle tree whose leaves are chain's batch hashes. If some chain on some settlement layer has exactly one batch, then its Merkle tree consists of exactly one node which is leaf and root at the same time. Merkle proof of this leaf is empty. However, function `MailboxFacet._proveL2LeafInclusion` forbids empty Merkle proof of batch leaf in two places: directly (`if (batchLeafProofLen == 0) { assume that root is stored here on L1 }`) and through `Merkle.calculateRootMemory` call.

Therefore, it will be impossible to prove logs sent in this batch until the second batch will be executed and sent to L1.



Also, malicious admin can execute one batch on settlement layer and then migrate chain back to L1. In such case, it will never be possible to prove logs sent in this batch.

**Recommendation:** In the first element of proof (where `logLeafProofLen`, `batchLeafProofLen` are encoded) encode also additional flag which distinguishes two scenarios — “Merkle proof is empty” or “root is stored here on L1”. Also do not revert in `Merkle._validatePathLengthForSingleProof` in case `_pathLength == 0`.

**ZKsync:** *Fixed* in the [dca7b558](#) commit.

**Audittens:** *Approved*.

## 6.2.6 H-06. NTV.assetId records can be overwritten

**Context:**

[l1-contracts/contracts/bridge/ntv/NativeTokenVault.sol:56,340](#)  
[l1-contracts/contracts/bridge/ntv/L2NativeTokenVault.sol:91](#)  
[l1-contracts/contracts/bridge/asset-router/L2AssetRouter.sol:155](#)

**Description:** Functions `NTV.registerToken` and `L2NTV.setLegacyTokenAssetId` allow to overwrite existing records of `assetId` mapping. It has the following impacts:

1. When some L2-native token (let ID1 be the corresponding asset id for it) is being withdrawn to L1 for the first time, bridged token T is being deployed on L1 and [NativeTokenVault.sol:421](#) stores `L1NTV.assetId[T] = ID1`. It allows to use `LEGACY_ENCODING_VERSION` to deposit asset ID1 to other chains (both native L2 and not), and use legacy recovery mechanism `claimFailedDeposit` for failed deposits of asset ID1 ([L1Nullifier.sol:651](#)). However, at any time anyone can call `NTV.registerToken`, which overwrites `L1NTV.assetId[T]` with some other ID2 (which is native for L1). After that, using `LEGACY_ENCODING_VERSION/claimFailedDeposit` for depositing/recovering of token T will operate with new asset ID2. It leads to:
  - In short time after overwriting `L1NTV.assetId[T]`, when regular users (A) try to call `claimFailedDeposit`, it'll revert due to zero `L1NTV.chainBalance[L2][ID2]`.
  - In short time after overwriting `L1NTV.assetId[T]`, when regular users (B) try to deposit token T, thinking that it still corresponds to asset ID1, it'll be deposited as ID2 and `L1NTV.chainBalance[L2][ID2]` will be increased.
  - After previous step, users (A) can successfully call `claimFailedDeposit`, since `L1NTV.chainBalance[L2][ID2]` is now increased by users (B).
  - After previous step, users (B) will realize that they deposited wrong asset ID2, but will not be able to withdraw it due to the decreased `L1NTV.chainBalance[L2][ID2]` by users (A). Therefore funds of users (B) remain permanently frozen, since all their withdrawals from L2 will be parsed on L1 ([L1Nullifier.sol:607,620](#)) as asset ID2 (not ID1, for which `chainBalance` exists).
2. On L2 when some native token T exists, it can be withdrawn using `L2AssetRouter.withdrawToken` function. However, at any time anyone can call `L2NTV.setLegacyTokenAssetId(T)`, which overwrites `L2NTV.assetId[T]` with non-existing `assetId` and therefore blocks use of `L2AssetRouter.withdrawToken` until `NTV.registerToken(T)` is manually called. By constantly calling `L2NTV.setLegacyTokenAssetId(T)`, griefers can prevent regular users from use of `L2AssetRouter.withdrawToken` function.

**Recommendation:** It's possible to fix the first impact of issue in two different ways depending on the project's needs:

1. If the functionality of depositing non-native assets from L1 to L2 through `LEGACY_ENCODING_VERSION` should exist, it's enough to forbid native registration of the bridged tokens ([NativeTokenVault.sol:340](#)) by adding the require statement `require(assetId[nativeToken] == 0);`. However, such an approach limits the functionality of using bridged tokens as separate native assets.
2. Otherwise, it's better to remove `assetId` from the NTV's storage and instead compute it everywhere as a native asset id for the current chain. Note, that it'll not break the backward compatibility, since the on-chain codebase doesn't support non-native tokens at all, and the currently reviewed codebase doesn't support depositing non-native assets through `LEGACY_ENCODING_VERSION` as soon as anyone calls `NTV.registerToken` (which can be done immediately after token's contract deployment).



To fix the second impact of the issue, the following has to be done:

1. `L2NTV.setLegacyTokenAssetId` function should check that it's being called when `L2_LEGACY_SHARED_BRIDGE` exists and stores information about `_l2TokenAddress`: `require(L2_LEGACY_SHARED_BRIDGE != address(0)); require(l1TokenAddress != address(0));`
2. `L2NTV.registerToken` function should check that `_nativeToken` doesn't correspond to the legacy bridged token, not yet registered by the `setLegacyTokenAssetId`: `require(L2_LEGACY_SHARED_BRIDGE == address(0) || L2_LEGACY_SHARED_BRIDGE.l1TokenAddress(_nativeToken) == address(0))`.

**ZKsync:** *Fixed* in the [7673edf9](#) and [2bcd5de9](#) commits.

**Audittens:** *Approved*.

## 6.2.7 H-07. Legacy withdrawals on L2 can be exploited to steal users' funds

**Context:**

[11-contracts/contracts/bridge/asset-router/L2AssetRouter.sol:284,297,307,316](#)

[11-contracts/contracts/bridge/L2SharedBridgeLegacy.sol:103-108](#)

**Description:** The current `L2SharedBridge` implements the `withdraw` function that successfully works only with `_l2Token` deployed by the bridge itself. However, the new implementation can successfully work with the malicious `_l2Token` as well, which can be exploited by the attacker.

**Attack scenario:** Imagine some passive vault on L2 that allows bridged tokens to be deposited into it and withdrawn to L1 (where withdrawals are being processed internally via `L2SharedBridge.withdraw`). In such a setup, the attacker can steal all the funds from the vault by taking the following steps:

1. Deploy a token `T` on L2 that implements the `l1Address` function that returns an address `A`. Here `A` is the address of some real token on L1 that was bridged to L2 and deposited to the vault above.
2. Mint a necessary amount of `T` for himself.
3. Deposit all minted amount of `T` into the vault.
4. Withdraw all deposited amount of `T` from the vault. During this step, the vault checks that the attacker indeed has deposited a corresponding amount of `T`, and calls `L2SharedBridgeLegacy.withdraw` with `_l2Token = T`. Since `T.l1Address() = A`, asset to withdraw corresponds to the token `A` ([L2AssetRouter.sol:307](#)). Therefore the vault withdraws all its balance of corresponding to token `A` asset directly to the attacker on L1.
5. Repeat steps 1-4 with all other tokens held by the vault.

**Recommendation:** Replace the implementation of the `L2AssetRouter.l1TokenAddress` with the following:

```
function l1TokenAddress(address _l2Token) public view returns (address) {
    require(L2_LEGACY_SHARED_BRIDGE != address(0));
    return IL2SharedBridgeLegacy(L2_LEGACY_SHARED_BRIDGE).l1TokenAddress(_l2Token);
}
```

**ZKsync:** *Fixed* in the [a9a28a2e](#) commit.

**Audittens:** *Approved*.



## 6.3 Medium severity findings

### 6.3.1 M-01. WETH withdrawing and recovering is impossible

**Context:**

[l1-contracts/contracts/bridge/ntv/NativeTokenVault.sol:92](#)

**Description:** Currently deployed L1SharedBridge forbids to deposit WETH token, but allows to withdraw and recover it for old deposits done through L1ERC20Bridge. When L1SharedBridge was introduced, ~ 17.9 WETH was transferred to it from L1ERC20Bridge. Since then, 19 successful withdrawals **happened**, resulting in ~ 8.8 WETH remaining on the L1SharedBridge balance.

Since [NativeTokenVault.sol:92](#) forbids to register WETH token, it not only forbids to deposit WETH token, but also forbids to withdraw and recover WETH for already existing deposits. Transferring liquidity of WETH from L1SharedBridge to L1NativeTokenVault is also impossible due to the necessity of token registration at [L1NativeTokenVault.sol:99](#).

Therefore introducing NativeTokenVault leads to the permanent freezing of all WETH liquidity in the L1SharedBridge.

**Recommendation:** Instead of forbidding WETH token registration at [NativeTokenVault.sol:92](#), forbid only WETH deposits by adding the corresponding restriction to the `_bridgeBurnNativeToken` function at [NativeTokenVault.sol:255](#).

**ZKsync:** *Fixed* in the [9db7c17a](#) commit.

**Audittens:** *Approved*.

### 6.3.2 M-02. Depositing through L1ERC20Bridge can spend funds twice

**Context:** [l1-contracts/contracts/bridge/L1ERC20Bridge.sol:206-210](#)

**Description:** [L1ERC20Bridge.sol:197](#) uses `L1_ASSET_ROUTER.depositLegacyErc20Bridge` to handle deposits. To calculate `_assetId`, [L1AssetRouter.sol:514](#) uses `_ensureTokenRegisteredWithNTV` function, which returns existing `assetId` from NTV if it's not zero. It allows to use L1ERC20Bridge for depositing non-native tokens.

However, when doing so, a user will spend funds twice:

1. At [L1ERC20Bridge.sol:229](#) when transferring funds to L1ERC20Bridge.
2. At [NativeTokenVault.sol:208](#) when burning funds by L1NTV.

For such a scenario, first half of the funds transferred to the L1ERC20Bridge will remain unused and permanently frozen there, since the approval to `L1_ASSET_ROUTER` is nullified at [L1ERC20Bridge.sol:206-210](#).

Additionally, if the transaction later fails on L2, `claimFailedDepositLegacyErc20Bridge` will not work as well due to incorrectly calculated `assetId` at [L1Nullifier.sol:708](#), assuming native token was deposited.

**Recommendation:** Since L1ERC20Bridge is not supposed to handle deposit operations with non-native tokens, change nullifying approval at [L1ERC20Bridge.sol:206-210](#) with the require statement that enforces all the allowance has been consumed: `require(token.allowance(address(this), L1_ASSET_ROUTER) == 0);`.

**ZKsync:** *Fixed* in the [59d11f7a](#) commit.

**Audittens:** *Approved*.





### 6.3.3 M-03. L2AssetRouter.\_ensureTokenRegisteredWithNTV always returns 0

**Context:**

[l1-contracts/contracts/bridge/asset-router/L2AssetRouter.sol:135-139,160](#)

**Description:** Function `L2AssetRouter._ensureTokenRegisteredWithNTV` returns `bytes32` `assetId`, but this function doesn't contain return statement nor assignment to `assetId`, so it always returns 0.

This function is used by `L2AssetRouter.withdrawToken` where its return value is used as asset id to be withdrawn. However, because of the mentioned bug, it will be equal to 0, so withdrawal will fail.

**Recommendation:** Add assignment `assetId = nativeTokenVault.assetId(_token)` at the end of the `L2AssetRouter._ensureTokenRegisteredWithNTV` function.

**ZKsync:** *Fixed* in the [7663cfc5](#) commit.

**Audittens:** *Approved.*

## 6.4 Low severity findings

### 6.4.1 L-01. Whitelisting of settlement layers affects MailboxFacet.\_proveL2LeafInclusion of all chains

**Context:**

[l1-contracts/contracts/state-transition/chain-deps/facets/Mailbox.sol:276](#)

[l1-contracts/contracts/bridgehub/Bridgehub.sol:299-308](#)

**Description:** `Bridgehub.registerSettlementLayer` is used to allow or disallow chains to settle on certain settlement layers. Also, `MailboxFacet._proveL2LeafInclusion` allows to prove leaf inclusion through any settlement layer whitelisted in Bridgehub. This means that `Bridgehub.registerSettlementLayer` has non-obvious consequences:

- when settlement layer is whitelisted, each chain is forced to trust this settlement layer (even chains which never settled on this settlement layer);
- when settlement layer is un-whitelisted, no chain can prove leaf inclusion through this settlement layer (even chains which settled on this settlement layer).

**Recommendation:** For each chain, store a mapping of used settlement layers, and allow proving only through them.

**ZKsync:** *Acknowledged.* Expected, since governance is expected to whitelist only chains with correct CTM, which assures correct tree building and so it is assumed that gateway can never produce malicious outputs (even if gateway admin is malicious). The proposed feature may be added in future versions if needed.

### 6.4.2 L-02. L1Nullifier.\_isPreSharedBridgeDepositOnEra behaviour is changed from on-chain version

**Context:**

[l1-contracts/contracts/bridge/L1Nullifier.sol:487-505](#)

[on-chain L1SharedBridge.sol:550](#)

**Description:** During upgrade, `L1Nullifier` contract will replace currently deployed `L1SharedBridge` (proxy, [implementation](#)). Function `L1SharedBridge._isEraLegacyDeposit` is renamed to `L1Nullifier._isPreSharedBridgeDepositOnEra`, its behaviour is kept the same except one check: “`_l2TxNumberInBatch < eraLegacyBridgeLastDepositTxNumber`” is replaced with “`_l2TxNumberInBatch <= eraLegacyBridgeLastDepositTxNumber`”. So there is one transaction which was treated as non-legacy but will be treated as legacy.

**Recommendation:** Depending on whether this transaction is legacy or not, either keep new implementation or restore the old one. In the latter case, it would be better to change variables' names and/or comments describing them accordingly.

**ZKsync:** *Fixed* in the [4c27dc67](#) commit.



Audittens: *Approved.*

#### 6.4.3 L-03. Allow to call AdminFacet.setPubdataPricingMode even after the first batch

Context:

[11-contracts/contracts/state-transition/chain-deps/facets/Admin.sol:129-132](#)

**Description:** Function AdminFacet.setPubdataPricingMode has a restriction that it can be called only when no batches are committed. Comment “Validium mode can be set only before the first batch is processed” makes sense. However, this function affects only pricing while data availability can be changed using function AdminFacet.setDAValidatorPair. Therefore, this restriction doesn’t protect chain from becoming validium. On the other hand, if changing data availability using setDAValidatorPair is allowed for some chains, setPubdataPricingMode should be also allowed to make pricing correct.

**Recommendation:** Remove this restriction.

**ZKsync:** *Fixed* in the [fd7e4cd6](#) commit.

Audittens: *Approved.*

#### 6.4.4 L-04. In RollupL1DAValidator, blob can be published long before it is really used

Context:

[da-contracts/contracts/RollupL1DAValidator.sol:22,148](#)

**Description:** RollupL1DAValidator allows to publish data using either calldata or blobs. Blobs are much cheaper because they are not completely available to EVM and are stored only for 4096 epochs (approximately 18 days). Also RollupL1DAValidator allows to publish blob in a separate transaction — anytime before batch is executed. This allows malicious validator to publish blob long before batch is executed (or even before chain is created) — at the time when batch is executed, this data won’t be available to users.

**Recommendation:** In RollupL1DAValidator, replace mapping(bytes32 blobCommitment => bool isPublished) public publishedBlobCommitments with mapping(bytes32 blobCommitment => uint256 publishBlockNumber) public publishedBlobCommitments.

In function publishBlobs, set publishedBlobCommitments[blobCommitment] = block.number.

In function checkDA, check require(block.number - publishedBlobCommitments[prepublishedCommitment] <= 4096 \* 32 / 2, "not published") so that one has some time (half of whole blob living time) to retrieve the blob.

**ZKsync:** *Fixed* in the [ae91ae43](#) commit.

Audittens: *Approved.*

#### 6.4.5 L-05. AdminFacet.forwardedBridgeRecoverFailedTransfer has redundant check

Context:

[11-contracts/contracts/state-transition/chain-deps/facets/Admin.sol:356](#)

**Description:** AdminFacet.forwardedBridgeRecoverFailedTransfer contains require(\_depositSender == s.admin, "Af: not chainAdmin") which effectively checks that admin isn’t changed during failed migration. This check doesn’t prevent any issues. Instead, it just creates ability to accidentally lose chain.

**Recommendation:** Remove this require statement.

**ZKsync:** *Fixed* in the [2c1aad9](#) commit.

Audittens: *Approved.*





#### 6.4.6 L-06. Bridging system only partially supports L2-value

**Context:**

[l1-contracts/contracts/bridgehub/IBridgehub.sol:26](#)  
[l1-contracts/contracts/bridgehub/Bridgehub.sol:505,540,555](#)  
[l1-contracts/contracts/bridge/asset-router/L1AssetRouter.sol:225,264](#)  
[l1-contracts/contracts/bridge/asset-router/L2AssetRouter.sol:122,126](#)  
[l1-contracts/contracts/bridge/asset-router/AssetRouterBase.sol:109,112,131,145](#)  
[l1-contracts/contracts/bridge/interfaces/IAssetHandler.sol:28](#)

**Description:** Bridging system on L1 is designed to support asset handlers which use L2-value:

- `Bridgehub.requestL2TransactionTwoBridges` takes parameter `_request.l2Value`, passes it to `IL1AssetRouter(_request.secondBridgeAddress).bridgehubDeposit` and later creates transaction with this L2-value;
- `L1AssetRouter.bridgehubDeposit` takes parameter `_value` and passes it to `IAssetHandler(l1AssetHandler).bridgeBurn` for verification;
- on L2, this L2-value is expected to be sent to `IAssetHandler.bridgeMint` which is payable.

However, such transaction will fail on L2 because `L2AssetRouter.finalizeDeposit` is not payable and doesn't pass `msg.value` to `IAssetHandler.bridgeMint` function calls.

**Recommendation:** Make function `L2AssetRouter.finalizeDeposit` payable and pass `msg.value` to `IAssetHandler.bridgeMint` function calls in function `AssetRouterBase._finalizeDeposit`.

**ZKsync:** *Fixed* in the [ed7da5dc](#) commit.

**Audittens:** *Approved*.

#### 6.4.7 L-07. PermanentRestriction.allowL2Admin is not resistant to CREATE2 collision

**Context:**

[l1-contracts/contracts/governance/PermanentRestriction.sol:100-120](#)  
[l1-contracts/contracts/governance/L2AdminFactory.sol:48](#)

**Description:** Function `PermanentRestriction.allowL2Admin` is used to permissionlessly register any `ChainAdmin` on L2 if its address is derived from approved `L2AdminFactory`.

However, this registration is not resistant to CREATE2 collision. User of `ChainAdmin` can find `salt1` and `salt2` such that `L2_ADMIN_FACTORY` with `salt1` deploys `ChainAdmin` to the same address as some malicious contract with `salt2` deploys another malicious contract. This requires computing approximately  $2^{81}$  hashes. In this case it's possible to register such an address in `PermanentRestriction.allowL2Admin` and do anything with the chain on L2, which should be restricted by `PermanentRestriction`.

**Recommendation:** In `L2AdminFactory.deployAdmin`, either:

- use `CREATE` instead of `CREATE2` and modify `PermanentRestriction.allowL2Admin` accordingly, or
- restrict `_salt` to be small (e.g. `bytes4`) and remove ability to use `_additionalRestrictions`.

**ZKsync:** *Fixed* in the [02d6fbb2](#) commit.

**Audittens:** *Approved*.



#### 6.4.8 L-08. L2WrappedBaseToken.initializeV2 fails for already deployed contract

**Context:**

[l1-contracts/contracts/bridge/L2WrappedBaseToken.sol:71](#)

[l2-contracts/contracts/bridge/L2WrappedBaseToken.sol:51](#) in old version

**Description:** Function `L2WrappedBaseToken.initializeV2` has to be called during upgrade process because new implementation introduces new variables `nativeTokenVault` and `baseTokenAssetId` and changes meaning of the old variable `l2Bridge`. However, this call will fail because of modifier `reinitializer(2)` since previous version was already initialized with `reinitializer(2)`.

**Recommendation:** Replace `reinitializer(2)` with `reinitializer(3)`.

**ZKsync:** *Fixed* in the [aeade7e2](#) commit.

**Audittens:** *Approved*.

#### 6.4.9 L-09. L1AssetRouter and L1Nullifier don't have some legacy functions for backward compatibility

**Context:**

[l1-contracts/contracts/bridgehub/Bridgehub.sol:898-901](#)

**Description:** During upgrade, `L1Nullifier` contract will replace currently deployed `L1SharedBridge`. On the other hand, `Bridgehub.sharedBridge` function (which returned address of `L1SharedBridge`) now returns address of `L1AssetRouter`. Therefore, contracts that used `L1SharedBridge` legacy interface may access either `L1Nullifier` (if they stored constant address) or `L1AssetRouter` (if they call `Bridgehub.sharedBridge` each time) and expect the same legacy interface. In fact, these new contracts implement almost all legacy functions, but:

- `L1AssetRouter` doesn't implement `depositHappened`;
- `L1Nullifier` doesn't implement `finalizeWithdrawal`, `L1_WETH.TOKEN` and bridging interface (`bridgehubDeposit`, `bridgehubConfirmL2Transaction`).

**Recommendation:** Implement all these functions.

**ZKsync:** *Partially fixed* in the [1982d826](#) commit. We decided to only add `finalizeWithdrawal` for now. Adding `bridgehubDeposit`-like methods would introduce too much complexity (e.g. in this case we would have to consider cases when users approve funds to `L1Nullifier` etc).

**Audittens:** *Approved*.

#### 6.4.10 L-10. Settlement layer's admin can break system invariant

**Context:**

[l1-contracts/contracts/state-transition/chain-deps/facets/Mailbox.sol:486-492,368-383](#)

**Description:** Settlement layers system doesn't support situations when there is a chain `C4` which settles on settlement layer `SL3` which itself settles on settlement layer `SL2` which settles on `L1`. For example, in such scenario, requesting priority transactions from `L1` to `C4` doesn't work because `MailboxFacet.requestL2TransactionToGatewayMailbox` of `SL3` doesn't call `MailboxFacet.requestL2TransactionToGatewayMailbox` of `SL2`.

However, the code doesn't explicitly restrict this situation — malicious admin of `SL3` can migrate it to `SL2` to put the system into an unsupported state. There are two possible (independent) attack scenarios:

1. In such a case, if anyone sends priority transaction from `L1` to `C4`, it will be stored in priority queue of `C4`'s diamond proxy on `L1` but not in priority queue of `C4`'s diamond proxy on `SL3`. Then admin of `SL3` can migrate it back to `L1`. After that, new priority transactions will be stored in priority queues of both `C4`'s diamond proxies. However, the first transaction will forever stay in priority queue of `C4`'s diamond proxy on `L1` and will not allow `C4` to be migrated back to `L1` because of this check: [Admin.sol:315-320](#).
2. If admin of `SL2` is malicious too, they can migrate `SL2` to `SL3` to create a loop where `SL3` settles on `SL2`, `SL2`



settles on SL3 and this structure is completely detached from L1. In such case, it will be impossible to migrate any SL to L1 to return back to the correct state because this would require sending message from other SL to L1 which is impossible since none of them settles on L1.

**Recommendation:** Forbid to migrate a chain to SL if it's whitelisted as settlement layer. Forbid to whitelist a chain as settlement layer if it's already on SL.

**ZKsync:** *Fixed* in the [dc549e25](#) commit.

**Audittens:** *Approved.*

#### 6.4.11 L-11. Operator can force user to execute the same transaction twice

**Note:** This finding is out of scope of this audit.

**Description:** Consider the following scenario:

1. User requests transaction directly to L2.
2. Operator ignores user.
3. User is forced to request the same transaction from L1.
4. Operator executes both transactions.

This way operator can force user to execute the same transaction twice (which means e.g. transferring twice more funds than expected and pay for transaction execution twice). User cannot avoid this because priority transaction doesn't increment user's nonce.

**Recommendation:** Add optional field "user's nonce" to priority transaction which has to be validated on L2 under usual rules of nonce management.

**ZKsync:** *Acknowledged.* It is a known issue with the current system.

## 6.5 Informational findings

### 6.5.1 I-01. ChainTypeManager.revertBatches fails unless ChainTypeManager is manually set as validator

**Context:**

[l1-contracts/contracts/state-transition/ChainTypeManager.sol:283](#)

[l1-contracts/contracts/state-transition/chain-deps/facets/Executor.sol:540](#)

**Description:** ChainTypeManager.revertBatches calls ExecutorFacet.revertBatchesSharedBridge which is callable only by validator because of modifier onlyValidator. So this call will fail unless ChainTypeManager is manually set as validator.

**Recommendation:** Apply modifier onlyValidatorOrChainTypeManager instead of onlyValidator to function ExecutorFacet.revertBatchesSharedBridge.

**ZKsync:** *Fixed* in the [fcc8de8](#) commit.

**Audittens:** *Approved.*

### 6.5.2 I-02. ChainTypeManager.registerSettlementLayer always fails

**Context:**

[l1-contracts/contracts/state-transition/ChainTypeManager.sol:443](#)

[l1-contracts/contracts/bridgehub/Bridgehub.sol:305](#)

**Description:** ChainTypeManager.registerSettlementLayer calls Bridgehub.registerSettlementLayer which is callable only by the owner because of modifier onlyOwner. So this call will always fail.



**Recommendation:** Either remove the function `ChainTypeManager.registerSettlementLayer` (since `Bridgehub.registerSettlementLayer` already provides the same functionality), or update function `Bridgehub.registerSettlementLayer` to allow calls both from owner and registered CTMs.

**ZKsync:** *Fixed* in the [f54121e6](#) commit.

**Audittens:** *Approved.*

### 6.5.3 I-03. MailboxFacet.proveL1ToL2TransactionStatusViaGateway function is empty

**Context:**

[11-contracts/contracts/state-transition/chain-deps/facets/Mailbox.sol:122-130](#)

**Description:** Function `MailboxFacet.proveL1ToL2TransactionStatusViaGateway` is empty.

**Recommendation:** Remove this function entirely since `MailboxFacet.proveL1ToL2TransactionStatus` already works for chains that settle on gateway.

**ZKsync:** *Fixed* in the [9016bd39](#) commit.

**Audittens:** *Approved.*

### 6.5.4 I-04. Wrong expected message length in `L1Nullifier._parseL2WithdrawalMessage`

**Context:**

[11-contracts/contracts/bridge/L1Nullifier.sol:617](#)

**Description:** In `L1Nullifier._parseL2WithdrawalMessage`, in case when `bytes4(functionSignature) == IAssetRouterBase.finalizeDeposit.selector`, it's checked that `_l2ToL1message.length >= 36`. However, actually message length has to be at least `bytes4 selector + uint256 originChainId + bytes32 assetId = 4 + 32 + 32 = 68`.

**Recommendation:** Replace 36 with 68.

**ZKsync:** *Fixed* in the [6bd1fb65](#) commit.

**Audittens:** *Approved.*

### 6.5.5 I-05. `NativeTokenVault.bridgeMint` emits event `BridgeMint` twice

**Context:**

[11-contracts/contracts/bridge/ntv/NativeTokenVault.sol:109-164](#)

**Description:** `NativeTokenVault.bridgeMint` always calls one of the internal functions: either `_bridgeMintNativeToken` or `_bridgeMintBridgedToken`. Both of these functions emit event `BridgeMint`. Then `NativeTokenVault.bridgeMint` itself emits exactly the same event again. Therefore, this event is anyway emitted twice.

**Recommendation:** Remove emit either in `NativeTokenVault.bridgeMint`, or in both `_bridgeMintNativeToken` and `_bridgeMintBridgedToken`.

**ZKsync:** *Fixed* in the [db410881](#) commit.

**Audittens:** *Approved.*



### 6.5.6 I-06. ERC20 getters of L2BaseToken always return information about Ether

**Context:**

[system-contracts/contracts/L2BaseToken.sol:131-147](#)

**Description:** ERC20 getters (name, symbol and decimals) of L2BaseToken always return information about Ether (Ether, ETH and 18). However, on chains where base token is not Ether, correct information about base token should be returned.

**Recommendation:** Information about the base token may be sent from L1 and set to the variables at genesis upgrade together with `setChainId`.

**ZKsync:** *Acknowledged*. It is a known issue and these methods have [instability warning comment](#) above them, so we may delete them in future releases.

### 6.5.7 I-07. IERC20.approve doesn't work with certain tokens

**Context:**

[l1-contracts/contracts/bridge/L1ERC20Bridge.sol:207,230](#)

**Description:** It's better to use `SafeERC20.forceApprove` instead of `IERC20.approve` because some tokens don't return `bool success`.

**Recommendation:** Replace `approve` with `forceApprove`.

**ZKsync:** *Fixed* in the [5c11f85a](#) commit.

**Audittens:** *Approved*.

### 6.5.8 I-08. AdminFacet.setPriorityTxMaxGasLimit and AdminFacet.setTokenMultiplier should be onlyL1

**Context:**

[l1-contracts/contracts/state-transition/chain-deps/facets/Admin.sol:83,114](#)

**Description:** Variables `s.priorityTxMaxGasLimit`, `s.baseTokenGasPriceMultiplierNominator` and `s.baseTokenGasPriceMultiplierDenominator` are used only on L1 (when priority transaction is requested), so it makes sense to set them only on L1.

**Recommendation:** Apply modifier `onlyL1` to functions `AdminFacet.setPriorityTxMaxGasLimit` and `AdminFacet.setTokenMultiplier`.

**ZKsync:** *Fixed* in the [f7364fb3](#) commit.

**Audittens:** *Approved*.

### 6.5.9 I-09. Unknown function selector is used in isTransactionAllowed

**Context:**

[l1-contracts/contracts/transactionFilterer/GatewayTransactionFilterer.sol:94](#)

**Description:** Function `GatewayTransactionFilterer.isTransactionAllowed` checks that transaction comes from `L1AssetRouter` and has correct function selector. Selectors `IL2AssetRouter.setAssetHandlerAddress` and `IAssetRouterBase.finalizeDeposit` are indeed functions of `L2AssetRouter`. However, the third checked selector `IL2Bridge.finalizeDeposit` is not implemented anywhere in the project. Therefore, it can be removed here.

**Recommendation:** Remove `"&& IL2Bridge.finalizeDeposit.selector != l2TxSelector"`.

**ZKsync:** *Fixed* in the [19262a20](#) commit.

**Audittens:** *Approved*.



### 6.5.10 I-10. Approve in L1AssetRouter.depositLegacyErc20Bridge is redundant

#### Context:

[l1-contracts/contracts/bridge/asset-router/L1AssetRouter.sol:515](#)

[l1-contracts/contracts/bridge/ntv/L1NativeTokenVault.sol:250-256](#)

**Description:** Function `L1AssetRouter.depositLegacyErc20Bridge` calls `forceApprove` to grant allowance to the `nativeTokenVault`. Later, function `L1NativeTokenVault._depositFunds` checks if it has approval from `L1AssetRouter` — if so, it will transfer funds from `L1AssetRouter` instead of the original caller.

However, this process is redundant because `L1AssetRouter` should never have funds. Instead, during legacy deposit, `L1AssetRouter` transfers funds from legacy bridge directly to `L1NativeTokenVault`.

**Recommendation:** Remove line [L1AssetRouter.sol:515](#) and function `L1NativeTokenVault._depositFunds`.

**ZKsync:** *Fixed* in the [3f2589c0](#) commit.

**Audittens:** *Approved*.

### 6.5.11 I-11. Several functions accept `msg.value` but don't use it

#### Context:

[l1-contracts/contracts/bridge/asset-router/L1AssetRouter.sol:229,225,242-249](#)

[l1-contracts/contracts/bridge/asset-router/L2AssetRouter.sol:97](#)

[l1-contracts/contracts/bridge/ntv/NativeTokenVault.sol:181,183,177,118](#)

[l1-contracts/contracts/bridge/ntv/L1NativeTokenVault.sol:196](#)

[l1-contracts/contracts/bridgehub/Bridgehub.sol:695,736,773](#)

#### Description:

1. [L1AssetRouter.sol:229,242-249](#) — `L1AssetRouter.bridgehubDeposit` is payable but “set asset handler counterpart” branch doesn't use `msg.value`;
2. [L1AssetRouter.sol:225,242-249](#) — in `L1AssetRouter.bridgehubDeposit` which accepts parameter `_value` (which is amount of base token to be sent to L2), “set asset handler counterpart” branch requests call to `L2AssetRouter.setAssetHandlerAddress` without checking that `_value` is non-zero but function `setAssetHandlerAddress` is non-payable ([L2AssetRouter.sol:97](#));
3. [NativeTokenVault.sol:181,183](#) — `NativeTokenVault.bridgeBurn` is payable but “bridged token” branch doesn't use `msg.value`;
4. [NativeTokenVault.sol:177](#), [Bridgehub.sol:691](#) — `NativeTokenVault.bridgeBurn` and `Bridgehub.bridgeBurn` accept parameter `_msgValue` (which is amount of base token to be sent to `bridgeMint` function on the other chain) without checking that it's non-zero but corresponding `NativeTokenVault.bridgeMint` and `Bridgehub.bridgeMint` don't use `msg.value`;
5. [NativeTokenVault.sol:118](#), [L1NativeTokenVault.sol:196](#), [Bridgehub.sol:695,736,773](#) — `NativeTokenVault.bridgeMint`, `L1NativeTokenVault.bridgeRecoverFailedTransfer`, `Bridgehub.bridgeBurn`, `Bridgehub.bridgeMint` and `Bridgehub.bridgeRecoverFailedTransfer` are payable but don't use `msg.value`.

In all these cases, if user accidentally sends `msg.value` / L2-value, they will lose it (with or without an ability to recover it).

#### Recommendation:

For items 1, 3 and 5, add statements `if (msg.value != 0) revert NonEmptyMsgValue()`.

For item 2, add statement `if (_value != 0) revert NonEmptyMsgValue()`.

For item 4, add statement `if (_msgValue != 0) revert NonEmptyMsgValue()`.

**ZKsync:** *Fixed* in the [2edde4c0](#) commit.

**Audittens:** *Approved*.



#### 6.5.12 I-12. `L2AssetRouter.withdraw` and `L2AssetRouter.withdrawToken` are marked as legacy

**Context:**

[l1-contracts/contracts/bridge/asset-router/L2AssetRouter.sol:141-162](#)

**Description:** `L2AssetRouter.withdraw` and `L2AssetRouter.withdrawToken` are new functions that will be introduced in the upcoming upgrade. However, they are placed under the comment that describes the block of “*legacy functions*”.

`L2AssetRouter.withdraw` is the only function in `L2AssetRouter` that can withdraw tokens. However, the comment above it says “*do not rely on this function, it will be deprecated in the future*”.

According to these comments, currently contracts and applications cannot withdraw assets without risk of deprecation in future.

**Recommendation:** Move the comment that describes the block of “*legacy functions*” under the `L2AssetRouter.withdraw` and `L2AssetRouter.withdrawToken` functions. Remove the comment “*do not rely on this function, it will be deprecated in the future*” above the `L2AssetRouter.withdraw` function.

**ZKsync:** *Fixed* in the [4324fe65](#) commit.

**Audittens:** *Approved.*

#### 6.5.13 I-13. `AccessControlRestriction.validateCall` may miss fallback role checks for invalid selectors

**Context:**

[l1-contracts/contracts/governance/AccessControlRestriction.sol:73](#)

**Description:** If the target contract doesn’t have a function matching the given selector (for example, this function is removed in an upgrade) the `AccessControlRestriction.validateCall` function still checks `requiredRoles` instead of `requiredRolesForFallback`. This allows callers who have roles for the specific selector but lack the required fallback role to access the fallback function.

**Recommendation:** Ensure the consistency between roles granted for some selectors and actual selectors of the functions of the target contracts. Specifically, ensure that roles associated with outdated selectors are removed.

**ZKsync:** *Acknowledged.*

#### 6.5.14 I-14. `L2AssetRouter.withdrawToken` supports both native tokens and tokens that came from another L2

**Context:**

[l1-contracts/contracts/bridge/asset-router/L2AssetRouter.sol:157-159](#)

**Description:** Function `L2AssetRouter.withdrawToken` reverts if token’s origin chain is L1. This means that it supports both native tokens (origin chain is current L2 or 0) and tokens from other L2s. However, it seems that this function is supposed to support only native tokens.

**Recommendation:** Replace condition “`recordedOriginChainId == L1_CHAIN_ID`” with “`recordedOriginChainId != 0 && recordedOriginChainId != block.chainid`”.

**ZKsync:** *Fixed* in the [201b7f05](#) commit.

**Audittens:** *Approved.*





#### 6.5.15 I-15. PermanentRestriction.\_validateRemoveRestriction blocks self-calls with data shorter than 4 bytes

**Context:**

[11-contracts/contracts/governance/PermanentRestriction.sol:203](#)

**Description:** Function `PermanentRestriction._validateRemoveRestriction` is designed to restrict calls to `removeRestriction` of the `ChainAdmin` who has this restriction. However, due to use of slice `_call.data[:4]`, this function will also revert when `_call.data` is shorter than 4 bytes. Currently this restricts only self-calls to `receive` but if `ChainAdmin` implementation will be extended in future, something else may become blocked too.

**Recommendation:** Before accessing slice, add one more check: “if (`_call.data.length < 4`) return;”.

**ZKsync:** *Fixed* in the [d335fe3b](#) commit.

**Audittens:** *Approved.*

#### 6.5.16 I-16. BridgeHelper.getERC20Getters doesn’t check whether calls are successful

**Context:**

[11-contracts/contracts/bridge/BridgeHelper.sol:29-31](#)

[11-contracts/contracts/bridge/BridgedStandardERC20.sol:117-127,135-140](#)

**Description:** When token is bridged for the first time, function `BridgeHelper.getERC20Getters` calls `name`, `symbol` and `decimals` functions of this token to be later decoded on the other chain. `BridgeHelper.getERC20Getters` doesn’t check success of these calls, so if some of them reverts, its revert data will be sent to the other chain. If this revert data will be accidentally parsed to the resulting type, bridged token’s getter will return incorrect value instead of revert.

**Recommendation:** In `BridgeHelper.getERC20Getters`, check whether getter reverted — if so, replace its revert data with empty bytes.

**ZKsync:** *Fixed* in the [39ca05d6](#) commit.

**Audittens:** *Approved.*

#### 6.5.17 I-17. GettersFacet.isFunctionFreezable and GettersFacet.isFacetFreezable are inconsistent

**Context:**

[11-contracts/contracts/state-transition/chain-deps/facets/Getters.sol:211-218,193-204](#)

**Description:** Function `GettersFacet.isFunctionFreezable` reverts when called with non-registered selector. Function `GettersFacet.isFacetFreezable` returns `false` when called with non-registered facet. For consistency, either both of them should revert, or both of them should return `false`.

**Recommendation:** In `GettersFacet.isFacetFreezable`, revert if condition “`selectorsArrayLen != 0`” doesn’t hold.

**ZKsync:** *Fixed* in the [4cab4c33](#) commit.

**Audittens:** *Approved.*

#### 6.5.18 I-18. New L3 → L1 message proof is not completely backward-compatible

**Context:**

[Nested L3 → L1 messages tree design for Gateway — Legacy support](#)  
[system-contracts/contracts/L1Messenger.sol:309](#)

**Description:** L3 → L1 messages were designed in a way to support legacy format of L2 → L1 logs proving: “*just provide a proof that assumes that stored `settledMessageRoot` is identical to local root, i.e. the hash of logs in the batch*”.





This statement is not a correct description of real proof format. In real proof format, the stored root is `fullRootHash = keccak256(localLogsRootHash || aggregatedRootHash)` — it may be seen as a root of non-standard tree where different leaves have different depths, but definitely not as “the hash of logs in the batch”.

**Recommendation:** If legacy support is necessary, send from `L1Messenger` and store in `s.l2LogsRootHashes` both roots — `fullRootHash` and `localLogsRootHash`. Otherwise, mention in documentation that there is only partial support of the legacy format.

**ZKsync:** *Acknowledged.* We will update documentation.

#### 6.5.19 I-19. Priority operation expiration timestamp is not stored in priority tree

**Context:**

[11-contracts/contracts/state-transition/chain-deps/facets/Mailbox.sol:572-583](#)

**Description:** When a priority operation is written to the priority tree, its expiration timestamp is ignored and not stored.

**Recommendation:** Store expiration timestamps corresponding to priority operations for use in future versions of the protocol.

**ZKsync:** *Acknowledged.* Expected, we will fix the issue when we’ll introduce full censorship resistance.

#### 6.5.20 I-20. Priority operation expiration should be increased for chains on settlement layer

**Context:**

[11-contracts/contracts/state-transition/chain-deps/facets/Mailbox.sol:480](#)

**Description:** `PRIORITY_EXPIRATION` is the amount of time the validator has to process the priority transaction. If the chain settles on some settlement layer, the amount of time needed to process the priority transaction depends not only on chain’s validator, but also on settlement layer’s validator:

1. settlement layer’s validator has `PRIORITY_EXPIRATION` time to execute `forwardTransactionOnGateway` transaction;
2. chain’s validator has `PRIORITY_EXPIRATION` time to execute the initial transaction on settlement layer;
3. if settlement layer’s validator ignores chain’s validator, then chain’s validator is forced to request a new transaction from L1 to settlement layer with execution of the batch which contains the initial transaction, so settlement layer’s validator has another `PRIORITY_EXPIRATION` time to execute this new transaction.

Therefore, priority operation expiration should be multiplied by 3.

**Recommendation:**

replace	<code>“_params.expirationTimestamp = uint64(block.timestamp +</code>
<code>PRIORITY_EXPIRATION)”</code>	<code>with “_params.expirationTimestamp = uint64(block.timestamp +</code>
<code>PRIORITY_EXPIRATION * (s.settlementLayer == address(0) ? 1 : 3))”.</code>	

**ZKsync:** *Acknowledged.* Expected, we will fix the issue when we’ll introduce full censorship resistance.

#### 6.5.21 I-21. In `ChainTypeManager._deployNewChain`, condition “`getZKChain(_chainId) != address(0)`” is never satisfied

**Context:**

[11-contracts/contracts/state-transition/ChainTypeManager.sol:355-358](#)

[11-contracts/contracts/bridgehub/Bridgehub.sol:363,370,752-753](#)

**Description:** `ChainTypeManager._deployNewChain` is called from two places:

- from `ChainTypeManager.createNewChain` which is called by `Bridgehub.createNewChain` where `_validateChainParams` checked that chain doesn’t exist ([Bridgehub.sol:363,370](#));



- from `ChainTypeManager.forwardedBridgeMint` which is called by `Bridgehub.bridgeMint` only if chain doesn't exist ([Bridgehub.sol:752-753](#)).

Therefore, condition “`getZKChain(_chainId) != address(0)`” in `ChainTypeManager._deployNewChain` cannot be satisfied.

**Recommendation:** Either completely remove this condition, or (to be safe) revert if it's satisfied.

**ZKsync:** *Fixed* in the [68d55a94](#) commit.

**Audittens:** *Approved*.

### 6.5.22 I-22. System design doesn't allow to register custom asset handlers for L2-native assets

**Context:**

[l1-contracts/contracts/bridge/asset-router/L1AssetRouter.sol:140,152,166,241](#)

[l1-contracts/contracts/bridge/asset-router/L2AssetRouter.sol:93,103](#)

**Description:** In `L1AssetRouter`, there are several functions that allow to use custom asset deployment trackers and asset handlers:

- `setAssetDeploymentTracker` — register deployment tracker for some asset on current chain;
- `setAssetHandlerAddressThisChain` — register handler for some asset which has deployment tracker on current chain;
- `_setAssetHandlerAddressOnCounterpart` in `bridgehubDeposit` — register deployment tracker on other chain for some asset which has deployment tracker on current chain.

`L2AssetRouter` can only:

- accept `_setAssetHandlerAddressOnCounterpart` requests from L1;
- `setAssetHandlerAddressThisChain` — register handler for some asset which has deployment tracker on current chain.

Therefore, it's impossible to create native asset on L2 (except `L2NativeTokenVault` which can be registered without `setAssetDeploymentTracker`).

**Recommendation:** Add missing functionality:

- `setAssetDeploymentTracker` on L2;
- `setAssetHandlerAddressOnCounterpart` on L2 and ability to receive such message on L1;
- `setAssetDeploymentTrackerAddressOnCounterpart` on L2 and ability to receive such message on L1 — this is needed to be able to bridge custom L2-native asset from L1 to another L2.

**ZKsync:** *Acknowledged*. Expected, it will be introduced in future releases.

### 6.5.23 I-23. Admin of a legacy chain can make `baseTokenAssetId` registered in Bridgehub

**Context:**

[l1-contracts/contracts/bridgehub/Bridgehub.sol:746-748](#)

**Description:** Function [Bridgehub.sol:732](#) assumes that `baseTokenAssetId` has been already registered in Bridgehub on L1. While this assumption is correct for new chains, for legacy chains it may not hold.

It allows the admins of legacy chains to register the corresponding `baseTokenAssetId` in Bridgehub on L1 by migrating their chain to a new settlement layer and back. Therefore the functionality of asset ids registration ([Bridgehub.sol:290](#)) is available not only to the owner or admin of Bridgehub, but also partially available to the admins of legacy chains.

**Recommendation:** Since `assetIdIsRegistered` is used only on L1, remove redundant registration at [Bridgehub.sol:748](#). Alternatively, if `assetIdIsRegistered` still has to be consistent, register the corresponding



baseTokenAssetId of legacy chains atomically inside setLegacyBaseTokenAssetId function. For this, the following line has to be added after [Bridgehub.sol:237](#): `assetIdIsRegistered[baseTokenAssetId[_chainId]] = true`.

**ZKsync:** *Fixed* in the [a71a0491](#) commit.

**Audittens:** *Approved*.

#### 6.5.24 I-24. Operator can force Executor to store systemLogs inconsistent with their hash

**Context:**

[11-contracts/contracts/state-transition/chain-deps/facets/Executor.sol:91-100,156-157,615](#)

**Description:** UnsafeBytes library is being used for processing systemLogs. However, in case emittedL2Logs.length % L2\_TO\_L1\_LOG\_SERIALIZE\_SIZE != 0, information about the last log will be partially read not from the systemLogs, but from the memory after it. In such a case, stored batch information at [Executor.sol:91-100](#) will not correspond to the hash at [Executor.sol:615](#) that is calculated for the batch commitment.

**Recommendation:** To guarantee consistency of system logs data with their hash and don't rely on the circuits' implementation, add the following line before processing systemLogs:  
`require(emittedL2Logs.length % L2_TO_L1_LOG_SERIALIZE_SIZE == 0).`

**ZKsync:** *Fixed* in the [924bf427](#) commit.

**Audittens:** *Approved*.

#### 6.5.25 I-25. Factory deps are not checked during the creation of new chain

**Context:**

[11-contracts/contracts/bridgehub/Bridgehub.sol:361,375](#)

[11-contracts/contracts/state-transition/ChainTypeManager.sol:416-420,426](#)

**Description:** During the creation of a new chain, CTM checks that the provided \_forceDeploymentData corresponds to the initialForceDeploymentHash. However, there's no similar check for the \_factoryDeps. It allows the chain's creator to pass the wrong \_factoryDeps, which will not allow execution of the genesis upgrade transaction. In such a case, the created chain will be unusable, and therefore other chain will have to be created, but with different not desired chainId.

**Recommendation:** Introduce new initialFactoryDepsHash variable and use it to validate the provided \_factoryDeps: `require(keccak256(abi.encode(_factoryDeps)) == initialFactoryDepsHash);`

**ZKsync:** *Acknowledged*, but we will not fix in this release.

#### 6.5.26 I-26. AssetRouterBase uses an incorrect storage gap size

**Context:**

[11-contracts/contracts/bridge/asset-router/AssetRouterBase.sol:53](#)

**Description:** According to the [OpenZeppelin upgradeable contracts standard](#), the size of the \_\_gap array is calculated so that the amount of storage slots used by a contract always adds up to the same number 50, but in the AssetRouterBase contract these add up to 49.

**Recommendation:** Change the \_\_gap in the AssetRouterBase contract to uint256[48].

**ZKsync:** *Fixed* in the [d947519c](#) commit.

**Audittens:** *Approved*.



### 6.5.27 I-27. Misleading comment in ChainTypeManager.registerSettlementLayer regarding settlement layer deployment

**Context:**

[11-contracts/contracts/state-transition/ChainTypeManager.sol:440-441](#)

**Description:** The comment at [ChainTypeManager.sol:440](#) states that it is required that the new settlement layer was deployed by the same ChainTypeManager, but the following check only ensures that it is registered in Bridgehub.

**Recommendation:** Ensure consistency between the mentioned comment and designed invariants of settlement layer registration.

**ZKsync:** *Fixed* in the [f54121e6](#) commit.

**Audittens:** *Approved.*

### 6.5.28 I-28. Misleading use of IBridgehub in abi.encodeCall

**Context:**

[11-contracts/contracts/state-transition/chain-deps/facets/Mailbox.sol:401](#)

**Description:** At [Mailbox.sol:401](#), IBridgehub(s.bridgehub).forwardTransactionOnGateway is used as the selector for the call, while the actual target of the call is L2\_BRIDGEHUB\_ADDR, not the s.bridgehub.

**Recommendation:** Replace IBridgehub(s.bridgehub).forwardTransactionOnGateway with the IBridgehub.forwardTransactionOnGateway.

**ZKsync:** *Fixed* in the [96e3a9b8](#) commit.

**Audittens:** *Approved.*

### 6.5.29 I-29. Misleading variable names corresponding to bridgehubData in Bridgehub contract

**Context:**

[11-contracts/contracts/bridgehub/Bridgehub.sol:709,753,782,716,759,789](#)

**Description:** It is confusing that bridgehubData.ctmData has different values in bridgeBurn, bridgeRecoverFailedTransfer and bridgeMint functions of the Bridgehub contract. The same happens with bridgehubData.chainData.

**Recommendation:** Rename the bridgehubData to bridgehubBurnData in bridgeBurn and bridgeRecoverFailedTransfer functions, and to bridgehubMintData in bridgeMint function.

**ZKsync:** *Fixed* in the [51363b9c](#) commit.

**Audittens:** *Approved.*

### 6.5.30 I-30. Redundant check for ASSET\_ROUTER in L1NativeTokenVault.receive

**Context:**

[11-contracts/contracts/bridge/ntv/L1NativeTokenVault.sol:72](#)

**Description:** At [L1NativeTokenVault.sol:72](#), the L1NativeTokenVault.receive function ensures that msg.sender is either L1\_NULLIFIER or ASSET\_ROUTER. But there is no way for msg.sender to be equal to ASSET\_ROUTER in this context, making the check for ASSET\_ROUTER redundant.

**Recommendation:** Remove the mentioned check.

**ZKsync:** *Fixed* in the [77ea5daa](#) commit.

**Audittens:** *Approved.*



### 6.5.31 I-31. Usage of raw keccak256 computation instead of `DataEncoding.encodeAssetId`

**Context:**

[l1-contracts/contracts/bridge/asset-router/L1AssetRouter.sol:144-146](#)

[l1-contracts/contracts/bridgehub/Bridgehub.sol:334](#)

[l1-contracts/contracts/common/libraries/DataEncoding.sol:65-67](#)

**Description:** At [L1AssetRouter.sol:144-146](#) and [Bridgehub.sol:334](#), the `assetId` is computed directly using `keccak256`. However, the function `DataEncoding.encodeAssetId` ([DataEncoding.sol:65-67](#)) is designed to provide such functionality.

**Recommendation:** Replace the direct computation with a call to `DataEncoding.encodeAssetId`.

**ZKsync:** *Fixed* in the [8409737f](#) commit. It was already partially fixed in a separate branch in the [7fb81ec2](#) commit.

**Audittens:** *Approved.*

### 6.5.32 I-32. Incorrect comment in `NativeTokenVault._bridgeBurnNativeToken` function regarding `msg.value` check

**Context:**

[l1-contracts/contracts/bridge/ntv/NativeTokenVault.sol:269-272](#)

**Description:** At [NativeTokenVault.sol:269-272](#), there is an enforcement that `msg.value` is not zero. However, the comment above it states “*The Bridgehub also checks this, but we want to be sure*”, while `Bridgehub` does not perform this check.

**Recommendation:** Update the comment to reflect the actual behavior.

**ZKsync:** *Fixed* in the [d994088c](#) commit.

**Audittens:** *Approved.*

### 6.5.33 I-33. Inconsistency between `MessageRoot` implementation and documentation

**Context:**

[Nested L3 → L1 messages tree design for Gateway](#)

[l1-contracts/contracts/bridgehub/MessageRoot.sol:116-118,141](#)

**Description:** The [documentation](#) states: “*In reality, `AggregatedRoot` will have a meaningful value only on `SyncLayer` and `L1`. On other chains it will be a root of an empty tree*”. In the `MessageRoot.getAggregatedRoot` function it is checked if `chainCount` is equal to zero, to return a root of an empty tree. However, at initialization of `MessageRoot`, `_addNewChain(block.chainid)` is called, therefore `chainCount` is greater than zero after initialization.

**Recommendation:** Replace the condition `chainCount == 0` with `chainCount == 1`.

**ZKsync:** *Acknowledged.* We will update documentation.

### 6.5.34 I-34. Incorrect implementation in the `L2AssetRouter`’s `onlyAssetRouterCounterpartOrSelf` modifier

**Context:**

[l1-contracts/contracts/bridge/asset-router/L2AssetRouter.sol:51-60](#)

**Description:** The implementation of the `L2AssetRouter`’s `onlyAssetRouterCounterpartOrSelf` modifier allows all calls where `_originChainId` is not equal to `L1_CHAIN_ID`.

**Recommendation:** Add an enforcement that the `_originChainId` is equal to `L1_CHAIN_ID`.

**ZKsync:** *Fixed* in a separate branch in the [21ac0837](#) commit.



Audittens: *Approved.*

#### 6.5.35 I-35. Inconsistent parameter naming of `_l1Asset` and `_l1Token`

Context:

[l1-contracts/contracts/bridge/L1Nullifier.sol:698](#)  
[l1-contracts/contracts/bridge/L1ERC20Bridge.sol:265](#)  
[l1-contracts/contracts/bridge/interfaces/IL1Nullifier.sol:44](#)  
[l1-contracts/contracts/bridge/asset-router/IL1AssetRouter.sol:36](#)

**Description:** There is inconsistency in the naming of the parameter representing the L1 token address across different contracts: in some places, it is referred to as `_l1Asset`, while in others it is called `_l1Token`.

**Recommendation:** Replace `_l1Asset` with `_l1Token` where L1 token is implied.

**ZKsync:** *Fixed* in the [b8970e5b](#) commit.

Audittens: *Approved.*

#### 6.5.36 I-36. Inconsistent use of `SUPPORTED_ENCODING_VERSION` constant in `BatchDecoder`

Context:

[l1-contracts/contracts/state-transition/libraries/BatchDecoder.sol:170](#)

**Description:** In [BatchDecoder.sol:170](#), `encodingVersion` is checked for equality with 0, while elsewhere ([BatchDecoder.sol:37,103](#)), the `SUPPORTED_ENCODING_VERSION` constant is used for the same check.

**Recommendation:** Replace the hardcoded 0 with `SUPPORTED_ENCODING_VERSION`.

**ZKsync:** *Fixed* in the [0be70a20](#) commit.

Audittens: *Approved.*

#### 6.5.37 I-37. Inconsistent variable naming for chain identifiers

Context:

[l1-contracts/contracts/bridge/ntv/NativeTokenVault.sol:133,153](#)  
[l1-contracts/contracts/bridge/asset-router/L2AssetRouter.sol:39,52,94](#)

**Description:** The parameter `_originChainId` is used inconsistently in the `_bridgeMintBridgedToken`, `_bridgeMintNativeToken` functions of the `NativeTokenVault` contract and in the `onlyAssetRouterCounterpart`, `onlyAssetRouterCounterpartOrSelf` modifiers and `setAssetHandlerAddress` function of the `L2AssetRouter` contract. In these functions and modifiers, it denotes the `chainId` from which the message originates, while in the rest of the codebase, `originChainId` is used to indicate the chain from which the token originated.

**Recommendation:** Replace `_originChainId` with `_chainId` in mentioned places.

**ZKsync:** *Fixed* in the [067b33ca](#) and [fe16f546](#) commits.

Audittens: *Approved.*

#### 6.5.38 I-38. Inconsistent naming of `assetAddress` and `assetHandlerAddress` in functions and events

Context:

[l1-contracts/contracts/bridge/asset-router/L2AssetRouter.sol:96](#)  
[l1-contracts/contracts/bridge/asset-router/IAssetRouterBase.sol:49](#)

**Description:** In most of the codebase, the naming `_assetHandlerAddress` is used to denote the address of an asset handler. However, there are two places in the codebase, where `_assetAddress` is used but the use of `_assetHandlerAddress` is appropriate:



- Function `setAssetHandlerAddress` takes `_assetAddress` as a parameter ([L2AssetRouter.sol:93-100](#)).
- The `AssetHandlerRegistered` event ([IAssetRouterBase.sol:49](#)) uses `_assetAddress` to denote the address of asset handler.

**Recommendation:** Rename `_assetAddress` to `_assetHandlerAddress` in the mentioned code.

**ZKsync:** *Fixed* in the [2a7723f1](#) commit. It was already partially fixed in a separate branch in the [21ac0837](#) commit.

**Audittens:** *Approved*.

### 6.5.39 I-39. Inconsistent behaviour of the legacy `L1AssetRouter.finalizeWithdrawal` function between legacy and new chains

**Context:**

[l1-contracts/contracts/bridge/asset-router/L1AssetRouter.sol:582](#)

**Description:** The function `L1AssetRouter.finalizeWithdrawal` is designed to support legacy withdrawals initiated by the legacy bridge. However, for new withdrawals that use new message format ([L2AssetRouter.sol:151,161,186-190](#)) and therefore corresponding withdrawal messages are sent by [L2AssetRouter.sol:189](#), the behaviour of this function with correctly provided parameters is different for different chains:

- For legacy chains, it will always revert, since `l2Sender` will be wrongly set as `legacyL2Bridge` ([L1AssetRouter.sol:582](#)).
- For new chains, it will succeed, since `l2Sender` will be correctly set as `L2_ASSET_ROUTER_ADDR`.

**Recommendation:** To have the same behaviour for all chains, allow to finalize only withdrawals sent by the legacy bridge: `require(legacyL2Bridge != address(0), {..., l2Sender: legacyL2Bridge, ...})`.

**ZKsync:** *Fixed* in the [235b42ea](#) commit.

**Audittens:** *Approved*.

### 6.5.40 I-40. Lack of reentrancy control in `L2AssetRouter` and `L2SharedBridgeLegacy`

**Context:**

[l1-contracts/contracts/bridge/asset-router/L2AssetRouter.sol:117,150,154,284,297](#)

[l1-contracts/contracts/bridge/L2SharedBridgeLegacy.sol:103,116](#)

**Description:** Functions of the `L2AssetRouter` and `L2SharedBridgeLegacy` contracts, that serve as entry points for deposit and withdrawal flows, are missing the `nonReentrant` modifier. This is inconsistent with the corresponding functionality on L1.

**Recommendation:** Add `nonReentrant` modifier to all of the entry points for finalizing deposits from L1 to L2 and withdrawing funds from L2 to L1.

**ZKsync:** *Fixed* in the [35b17442](#) commit.

**Audittens:** *Approved*.

